

Homework 4

Brandon Amaral, Monte Davityan, Nicholas Lombardo, Hongkai Lu

2022-10-03

1 Murphy Chapter 13 & ISLR Chapter 10 Summaries

Section 10.6 When to Use Deep Learning

The text shows three ways to determine when to use deep learning: 1: Occam's razor principle: when faced with several methods that give roughly equivalent performance, pick the simplest. 2: According to the text, "if we can produce models with the simpler tools that perform as well, they are likely to be easier to fit and understand, and potentially less fragile than the more complex approaches." 3: When the sample size of the data is extremely large and the interpretability of the model is not our main priority, we should use Deep Learning to be our first choice.

Section 10.7 Fitting a Neural Network

When fitting a neural network model, the text provides two general strategies: 1. Slow Learning 2. Regularization

10.7.1 Backpropagation Backpropagation is a process that the act of differentiation assigns a fraction of the residual to each of the parameters via the chain rule.

10.7.2 Regularization and Stochastic Gradient Descent Stochastic Gradient Descent is a process that we sample a small fraction of n observations' summation of the derivatives of $R_i(\theta)$ with respect to β_k and w_{kj} each time we compute a gradient step if n is large. In this process, Regularization is essential to avoid overfitting.

10.7.3 Dropout Learning Dropout learning is a way to randomly remove a fraction ϕ of the units in a layer when fitting the model.

10.7.4 Network Tuning There are three ways provided for network tuning in the text, which are "The number of hidden layers, and the number of units per layer", "The number of hidden layers, and the number of units per layer", and "The number of hidden layers, and the number of units per layer".

Section 10.8 Interpolation and Double Descent

Double descent is a phenomenon that the test error has a U-shape before the interpolation threshold is reached, then it descends again as an increasingly flexible model is fit. That may be applied to deep learning as well. However, even though double descent can occasionally occur in neural networks model, we do not rely on this behavior typically; meanwhile, it is important to know that the bias-variance trade-off always holds.

2 Reproduction of ISLR Lab 10.9

10.9.1 A Single Layer Network on the Hitters Data

First, We load the Hitters data and split it to training and test set.

```
Gitters <- na.omit(Hitters)
n <- nrow(Gitters)
set.seed(13)
ntest <- trunc(n/3)
testid <- sample(1:n, ntest)
```

Then, we fit a linear model and obtain the mean absolute prediction error, which is 254.6687.

```
lfit <- lm(Salary ~., data=Gitters[-testid,])
lpred <- predict(lfit, Gitters[testid,])
with(Gitters[testid,], mean(abs(lpred - Salary)))
```

```
## [1] 254.6687
```

Next, we fit a lasso model. Similarly, we calculate the absolute prediction error of the lasso model, which is slightly less than the result of linear model.

```
library(glmnet)
x <- scale(model.matrix(Salary ~. -1, data=Gitters))
y <- Gitters$Salary


cvfit <- cv.glmnet(x[-testid, ], y[-testid], type.measure='mae')
cpred <- predict(cvfit, x[testid,], s='lambda.min')
mean(abs(y[testid] - cpred))
```

```
## [1] 252.2994
```

Last, we fit the neural network with library keras. As the text written, it is a challenge to get keras running on computer. The error I got is “CondaSSLError: Encountered an SSL error. Most likely a certificate verification issue.” Still figure out a way to get this installing.

```
modnn <- keras_model_sequential() %>%
  layer_dense(units=50, activation='relu', input_shape = ncol(x)) %>%
  layer_dropout(rate=0.4) %>%
  layer_dense(units=1)

modnn %>% compile(loss='mse', optimizer=optimizer_rmsprop(), metrics=list('mean_absolute_error'))
history <- modnn %>% fit(
  x[-testid, ], y[-testid], epochs=200, batch_size=32,
  validation_data=list(x[testid,], y[testid])
)
plot(history, smooth = FALSE)
```



```
npred <- predict(modnn, x[testid,])
mean(abs(y[testid] - npred))
```

```
## [1] 351.8675
```

10.9.2 A Multilayer Network on the MNIST Digit Data

```
mnist <- dataset_mnist()
x_train <- mnist$train$x
g_train <- mnist$train$y
x_test <- mnist$test$x
g_test <- mnist$test$y
dim(x_train)
```

```
## [1] 60000    28    28
```

```
x_train <- array_reshape(x_train, c(nrow(x_train), 784))
x_test <- array_reshape(x_test, c(nrow(x_test), 784))
y_train <- to_categorical(g_train, 10)
y_test <- to_categorical(g_test, 10)
```

```
x_train <- x_train/255
x_test <- x_test/255
```

```
modelnn <- keras_model_sequential()
modelnn %>%
  layer_dense(units=256, activation="relu", input_shape=c(784)) %>%
```

```

layer_dropout(rate=0.4) %>%
layer_dense(units=128, activation='relu') %>%
layer_dropout(rate=0.3) %>%
layer_dense(units=10, activation='softmax')

#summary(modelnn)

modelnn %>% compile(loss='categorical_crossentropy', optimizer=optimizer_rmsprop(), metrics=c("accuracy"))

system.time(
  history <- modelnn %>%
    fit(x_train, y_train, epochs=30, batch_size=128, validation_split=0.2)
)

##      user      system elapsed
## 109.581    7.724    48.001

accuracy <- function(pred, truth) {
  mean(drop(as.numeric(pred)) == drop(truth)) }
modelnn %>% predict(x_test) %>% k_argmax() %>% accuracy(g_test)

## [1] 0.9814

modellr <- keras_model_sequential() %>%
  layer_dense(input_shape = 784, units = 10,
    activation = "softmax")
#summary(modellr)

modellr %>% compile(loss = "categorical_crossentropy",
  optimizer = optimizer_rmsprop(), metrics = c("accuracy"))
modellr %>% fit(x_train, y_train, epochs = 30,
  batch_size = 128, validation_split = 0.2)
modellr %>% predict(x_test) %>% k_argmax() %>% accuracy(g_test)

## [1] 0.9285

```

10.9.3 Convolutional Neural Networks

```

cifar100 <- dataset_cifar100()
names(cifar100)

## [1] "train" "test"

x_train <- cifar100$train$x
g_train <- cifar100$train$y
x_test <- cifar100$test$x
g_test <- cifar100$test$y
dim(x_train)

## [1] 50000    32    32     3

range(x_train[1,,, 1])

## [1] 13 255


x_train <- x_train / 255
x_test <- x_test / 255
y_train <- to_categorical(g_train, 100)

```

```
dim(y_train)
```

```
## [1] 50000 100
```

```
library(jpeg)
par(mar = c(0, 0, 0, 0), mfrow = c(5, 5))
index <- sample(seq(50000), 25)
for (i in index) plot(as.raster(x_train[i,,, ]))
```



```
model <- keras_model_sequential() %>%
  layer_conv_2d(filters = 32, kernel_size = c(3, 3),
    padding = "same", activation = "relu",
    input_shape = c(32, 32, 3)) %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_conv_2d(filters = 64, kernel_size = c(3, 3),
    padding = "same", activation = "relu") %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_conv_2d(filters = 128, kernel_size = c(3, 3),
    padding = "same", activation = "relu") %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_conv_2d(filters = 256, kernel_size = c(3, 3),
    padding = "same", activation = "relu") %>%
  layer_max_pooling_2d(pool_size = c(2, 2)) %>%
  layer_flatten() %>%
  layer_dropout(rate = 0.5) %>%
  layer_dense(units = 512, activation = "relu") %>%
  layer_dense(units = 100, activation = "softmax")
```

```
#summary(model)

model %>% compile(loss = "categorical_crossentropy",
  optimizer = optimizer_rmsprop(), metrics = c("accuracy"))
#history <- model %>% fit(x_train, y_train, epochs = 30,
history <- model %>% fit(x_train, y_train, epochs = 10,
  batch_size = 128, validation_split = 0.2)
model %>% predict(x_test) %>% k_argmax() %>% accuracy(g_test)

## [1] 0.4058
```

10.9.4 Using Pretrained CNN Models

```
img_dir <- "book_images"
image_names <- list.files(img_dir)
num_images <- length(image_names)
x <- array(dim = c(num_images, 224, 224, 3))
for (i in 1:num_images) {
  img_path <- paste(img_dir, image_names[i], sep = "/")
  img <- image_load(img_path, target_size = c(224, 224))
  x[i,, ] <- image_to_array(img)
}
x <- imagenet_preprocess_input(x)
```

```
model <- application_resnet50(weights = "imagenet")
#summary(model)
```

```
pred6 <- model %>% predict(x) %>%
  imagenet_decode_predictions(top = 3)
names(pred6) <- image_names
print(pred6)
```

```
## $flamingo.jpg
##   class_name class_description      score
## 1  n02007558      flamingo 0.926349819
## 2  n02006656      spoonbill 0.071699291
## 3  n02002556      white_stork 0.001228209
##
## $hawk_cropped.jpeg
##   class_name class_description      score
## 1  n01608432          kite 0.72270876
## 2  n01622779      great_grey_owl 0.08182590
## 3  n01532829      house_finch 0.04218867
##
## $hawk.jpg
##   class_name class_description      score
## 1  n03388043      fountain 0.2788654
## 2  n03532672          hook 0.1785550
## 3  n03804744          nail 0.1080727
##
## $huey.jpg
##   class_name      class_description      score
## 1  n02097474      Tibetan_terrier 0.50929713
## 2  n02098413          Lhasa 0.42209816
```

```
## 3 n02098105 soft-coated_wheaten_terrier 0.01695861
##
## $kitty.jpg
## class_name class_description score
## 1 n02105641 Old_English_sheepdog 0.83265942
## 2 n02086240 Shih-Tzu 0.04513892
## 3 n03223299 doormat 0.03299790
##
## $weaver.jpg
## class_name class_description score
## 1 n01843065 jacamar 0.49795377
## 2 n01818515 macaw 0.22193290
## 3 n02494079 squirrel_monkey 0.04287881
```

10.9.5 IMDb Document Classification

```
max_features <- 10000
imdb <- dataset_imdb(num_words = max_features)
c(c(x_train, y_train), c(x_test, y_test)) %<-% imdb
```

```
x_train[[1]][1:12]
```

```
## [1] 1 14 22 16 43 530 973 1622 1385 65 458 4468
```

```
word_index <- dataset_imdb_word_index()
decode_review <- function(text, word_index) {
  word <- names(word_index)
  idx <- unlist(word_index, use.names = FALSE)
  word <- c("<PAD>", "<START>", "<UNK>", "<UNUSED>", word)
  idx <- c(0:3, idx + 3)
  words <- word[match(text, idx, 2)]
  paste(words, collapse = " ")
}
decode_review(x_train[[1]][1:12], word_index)
```

```
## [1] "<START> this film was just brilliant casting location scenery story direction everyone's"
```

```
library(Matrix)
one_hot <- function(sequences, dimension) {
  seqlen <- sapply(sequences, length)
  n <- length(seqlen)
  rowind <- rep(1:n, seqlen)
  colind <- unlist(sequences)
  sparseMatrix(i = rowind, j = colind,
    dims = c(n, dimension))
}
```

```
x_train_1h <- one_hot(x_train, 10000)
x_test_1h <- one_hot(x_test, 10000)
dim(x_train_1h)
```

```
## [1] 25000 10000
```

```
nnzero(x_train_1h) / (25000 * 10000)
```

```
## [1] 0.01316987
```


```

set.seed(3)
ival <- sample(seq(along = y_train), 2000)

library(glmnet)
fitlm <- glmnet(x_train_1h[-ival, ], y_train[-ival],
  family = "binomial", standardize = FALSE)
classlmv <- predict(fitlm, x_train_1h[ival, ]) > 0
acclmv <- apply(classlmv, 2, accuracy, y_train[ival] > 0)

par(mar = c(4, 4, 4, 4), mfrow = c(1, 1))
plot(-log(fitlm$lambda), acclmv)

```



```

model <- keras_model_sequential() %>%
  layer_dense(units = 16, activation = "relu",
    input_shape = c(10000)) %>%
  layer_dense(units = 16, activation = "relu") %>%
  layer_dense(units = 1, activation = "sigmoid")
model %>% compile(optimizer = "rmsprop",
  loss = "binary_crossentropy", metrics = c("accuracy"))
history <- model %>% fit(x_train_1h[-ival, ], y_train[-ival],
  epochs = 20, batch_size = 512,
  validation_data = list(x_train_1h[ival, ], y_train[ival]))

history <- model %>% fit(
  x_train_1h[-ival, ], y_train[-ival], epochs = 20,
  batch_size = 512, validation_data = list(x_test_1h, y_test)
)

```


10.9.6 Recurrent Neural Networks

```
wc <- sapply(x_train, length)
median(wc)
```

Sequential Models for Document Classification

```
## [1] 178
```

```
sum(wc <= 500) / length(wc)
```

```
## [1] 0.91568
```

```
maxlen <- 500
x_train <- pad_sequences(x_train, maxlen = maxlen)
x_test <- pad_sequences(x_test, maxlen = maxlen)
dim(x_train)
```

```
## [1] 25000 500
```

```
dim(x_test)
```


```
## [1] 25000 500
```

```
x_train[1, 490:500]
```

```
## [1] 16 4472 113 103 32 15 16 5345 19 178 32
```

```
model <- keras_model_sequential() %>%
  layer_embedding(input_dim = 10000, output_dim = 32) %>%
  layer_lstm(units = 32) %>%
  layer_dense(units = 1, activation = "sigmoid")
```

```
model %>% compile(optimizer = "rmsprop",
  loss = "binary_crossentropy", metrics = c("acc"))
#history <- model %>% fit(x_train, y_train, epochs = 10,
history <- model %>% fit(x_train, y_train, epochs = 3,
  batch_size = 128, validation_data = list(x_test, y_test))
plot(history, smooth = FALSE)
```



```
xdata <- data.matrix(
  NYSE[, c("DJ_return", "log_volume", "log_volatility")]
)
istrain <- NYSE[, "train"]
xdata <- scale(xdata)
```

```
lagm <- function(x, k = 1) {
  n <- nrow(x)
  pad <- matrix(NA, k, ncol(x))
  rbind(pad, x[1:(n - k), ])
}
```

```
arframe <- data.frame(log_volume = xdata[, "log_volume"],
  L1 = lagm(xdata, 1), L2 = lagm(xdata, 2),
  L3 = lagm(xdata, 3), L4 = lagm(xdata, 4),
  L5 = lagm(xdata, 5)
)
```

```
arframe <- arframe[-(1:5), ]
istrain <- istrain[-(1:5)]
```

```
arfit <- lm(log_volume ~ ., data = arframe[istrain, ])
```

```

arpred <- predict(arfit, arframe[!istrain, ])
V0 <- var(arframe[!istrain, "log_volume"])
1 - mean((arpred - arframe[!istrain, "log_volume"])^2) / V0

```

Time Series Prediction

```
## [1] 0.413223
```

```

arframed <-
  data.frame(day = NYSE[-(1:5), "day_of_week"], arframe)
arfitd <- lm(log_volume ~ ., data = arframed[istrain, ])
arpredd <- predict(arfitd, arframed[!istrain, ])
1 - mean((arpredd - arframe[!istrain, "log_volume"])^2) / V0

```

```
## [1] 0.4598616
```

```

n <- nrow(arframe)
xrnn <- data.matrix(arframe[, -1])
xrnn <- array(xrnn, c(n, 3, 5))
xrnn <- xrnn[, , 5:1]
xrnn <- aperm(xrnn, c(1, 3, 2))
dim(xrnn)

```

```
## [1] 6046    5    3
```

```

model <- keras_model_sequential() %>%
  layer_simple_rnn(units = 12,
    input_shape = list(5, 3),
    dropout = 0.1, recurrent_dropout = 0.1) %>%
  layer_dense(units = 1)
model %>% compile(optimizer = optimizer_rmsprop(),
  loss = "mse")

```

```

history <- model %>% fit(
  xrnn[istrain, , ], arframe[istrain, "log_volume"],
  #   batch_size = 64, epochs = 200,
  batch_size = 64, epochs = 75,
  validation_data =
    list(xrnn[!istrain, , ], arframe[!istrain, "log_volume"])
)
kpred <- predict(model, xrnn[!istrain, , ])
1 - mean((kpred - arframe[!istrain, "log_volume"])^2) / V0

```

```
## [1] 0.407035
```

```

model <- keras_model_sequential() %>%
  layer_flatten(input_shape = c(5, 3)) %>%
  layer_dense(units = 1)

```

```

x <- model.matrix(log_volume ~ . - 1, data = arframed)
colnames(x)


```

```

## [1] "dayfri"      "daymon"      "daythur"
## [4] "daytues"     "daywed"      "L1.DJ_return"
## [7] "L1.log_volume" "L1.log_volatility" "L2.DJ_return"
## [10] "L2.log_volume" "L2.log_volatility" "L3.DJ_return"
## [13] "L3.log_volume" "L3.log_volatility" "L4.DJ_return"
## [16] "L4.log_volume" "L4.log_volatility" "L5.DJ_return"

```

```
## [19] "L5.log_volume"      "L5.log_volatility"
arnnd <- keras_model_sequential() %>%
  layer_dense(units = 32, activation = 'relu',
    input_shape = ncol(x)) %>%
  layer_dropout(rate = 0.5) %>%
  layer_dense(units = 1)
arnnd %>% compile(loss = "mse",
  optimizer = optimizer_rmsprop())
history <- arnnd %>% fit(
  # x[istrain, ], arframe[istrain, "log_volume"], epochs = 100,
  x[istrain, ], arframe[istrain, "log_volume"], epochs = 30,
  batch_size = 32, validation_data =
    list(x[!istrain, ], arframe[!istrain, "log_volume"])
)
plot(history, smooth = FALSE)
```



```
npred <- predict(arnnd, x[!istrain, ])
1 - mean((arframe[!istrain, "log_volume"] - npred)^2) / V0
```

```
## [1] 0.4698843
```

3 ISLR Problems

10.7

We fit a neural network with a single hidden layer and 10 units to the `Default` data, and compare its performance to a logistic regression model. First, we scale our data in a model matrix, excluding the usual

first column of 1's for an intercept term, and create a separate response vector of 1's and 0's.

```
x <- model.matrix(default ~ . -1, data = Default) |> scale()
y <- ifelse(Default$default == "Yes",1,0)
```

Next, we partition our data into training and test datasets, training on a random set of 80% of the data, and testing on the remaining 20%.

```
n <- nrow(x)
test_ind <- sample(1:n, n/5)

x_train <- x[-test_ind,]
x_test <- x[test_ind,]
y_train <- y[-test_ind]
y_test <- y[test_ind]
```

Now, we build the model architecture, defining a single layer model with 10 units in the hidden layer, with the usual ReLU activation function and a dropout rate of 20% for regularization. Since we are predicting a binary response, we use the sigmoid activation function in the output layer. Then, we compile the model, using the binary cross entropy loss function as a measure of fit.


```
nnmodel <- keras_model_sequential() |>
  layer_dense(units = 10, activation = "relu", input_shape = ncol(x)) |>
  layer_dropout(rate = 0.20) |>
  layer_dense(units = 1, activation = "sigmoid")

nnmodel |> compile(
  loss = "binary_crossentropy",
  optimizer = optimizer_rmsprop(),
  metrics = "accuracy")
```

We fit the model using a batch size of 30 with 40 epochs. We also further split the data into a third validation set, again with 20%, so we fit the model on 6400 training observations. Then, each epoch is $6400/30 \approx 214$ stochastic gradient descent steps. We plot the loss of the model alongside its

```
history <- nnmodel |>
  fit(x_train, y_train, epochs = 40, batch_size = 30, validation_split = 0.2)

plot(history, smooth = FALSE)
```



Now, we make our predictions on the test set and find the model's misclassification rate.

```
y_test_pred <- nnmodel |> predict(x_test)
NN_misclass_rate <- 1/length(y_test)*sum(y_test != round(y_test_pred))
```

Now, we compare this with a logistic regression model.

```
log_model <- glm(y_train ~ x_train, family = binomial(link = logit))
y_log_test_pred <- round(predict(log_model, newdata = data.frame(x_test), type = "response"))
logistic_misclass_rate <- 1/length(y_test)*sum(y_test != round(y_log_test_pred))
```

The table shows the misclassification rates for both models

```
misclass <- data.frame(Logistic = logistic_misclass_rate, NeuralNet = NN_misclass_rate)
knitr::kable(misclass)
```

Logistic	NeuralNet
0.1915	0.029

In the table, we see that the logistic regression model performed considerably worse, with a misclassification rate of about 19.15% versus the neural net's misclassification rate of about 2.9%.

10.8

```
img_dir <- "MyAnimalPics"
image_names <- list.files(img_dir)
num_images <- length(image_names)
```

```
x <- array ( dim = c(num_images , 224, 224, 3))
for (i in 1:num_images) {
  img_path <- paste (img_dir , image_names[i], sep = "/")
  img <- image_load (img_path, target_size = c(224, 224))
  x[i,,, ] <- image_to_array(img)
}
x <- imagenet_preprocess_input (x)
```

```
model <- application_resnet50(weights = "imagenet")
#summary (model)
```

```
pred6 <- model %>% predict (x) %>%
  imagenet_decode_predictions (top = 5)
names (pred6) <- image_names
print (pred6)
```

```
## $Bear.jpg
##   class_name class_description      score
## 1  n02363005          beaver 0.63836390
## 2  n02077923          sea_lion 0.15228654
## 3  n02504458 African_elephant 0.04073466
## 4  n02444819          otter 0.02228249
## 5  n02396427          wild_boar 0.01866793
##
## $Cat.jpg
##   class_name class_description      score
## 1  n03786901          mortar 0.36851746
## 2  n03794056      mousetrap 0.20677969
## 3  n04399382          teddy 0.05336932
## 4  n02328150          Angora 0.04163421
## 5  n02950826          cannon 0.03490378
##
## $Coyote.jpg
##   class_name class_description      score
## 1  n04275548      spider_web 0.738270223
## 2  n02325366      wood_rabbit 0.097948611
## 3  n02494079 squirrel_monkey 0.020395182
## 4  n02457408 three-toed_sloth 0.018195676
## 5  n02486410          baboon 0.008232873
##
## $Dog.jpg
##   class_name class_description      score
## 1  n02106662 German_shepherd 0.846454084
## 2  n02105412          kelpie 0.093938135
## 3  n02105162      malinois 0.021421773
## 4  n02106550      Rottweiler 0.007880141
## 5  n02099712 Labrador_retriever 0.005973050
##
## $Fish.jpg
##   class_name class_description      score
## 1  n03908714 pencil_sharpener 0.30948049
## 2  n03127747      crash_helmet 0.06694938
## 3  n04311174          steel_drum 0.05349804
## 4  n02843684          birdhouse 0.04055086
```

```

## 5  n03443371          goblet 0.04000980
##
## $Hamster.jpg
##   class_name class_description      score
## 1  n03793489          mouse 0.34858209
## 2  n02233338        cockroach 0.06264669
## 3  n02909870          bucket 0.04642543
## 4  n03794056        mousetrap 0.04376448
## 5  n04192698          shield 0.03958284
##
## $Lizard.jpg
##   class_name class_description      score
## 1  n13044778        earthstar 0.16892505
## 2  n07749582          lemon 0.09082820
## 3  n13054560          bolete 0.07577771
## 4  n07754684        jackfruit 0.04505501
## 5  n02219486          ant 0.04071513
##
## $Monkey.jpg
##   class_name class_description      score
## 1  n02489166  proboscis_monkey 0.18063334
## 2  n02493509          titi 0.14519861
## 3  n02492035        capuchin 0.14259827
## 4  n02492660    howler_monkey 0.10484368
## 5  n03017168          chime 0.09478551
##
## $Parrot.jpg
##   class_name class_description      score
## 1  n03991062          pot 0.32824290
## 2  n01580077          jay 0.22052953
## 3  n01882714          koala 0.11813702
## 4  n01818515          macaw 0.03252555
## 5  n02843684    birdhouse 0.02874277
##
## $Penguin.jpg
##   class_name class_description      score
## 1  n04523525          vault 0.30266422
## 2  n04553703    washbasin 0.10002303
## 3  n04005630          prison 0.06567439
## 4  n04040759          radiator 0.04754867
## 5  n03874599          padlock 0.02866268
##
## $Stingray.jpg
##   class_name class_description      score
## 1  n01496331    electric_ray 0.817473173
## 2  n01498041        stingray 0.154839367
## 3  n01924916        flatworm 0.020372352
## 4  n07734744        mushroom 0.002621449
## 5  n01950731        sea_slug 0.002227621

```

Interestingly, the model does decently well for semi- obvious, clear pictures such as the Stingray, monkey, and dog images. However, it does really poorly when the animal is not obvious such as the coyote, parrot, and fish.